



# Randomized algorithms 4

## Distributed computing

CS240

Spring 2020

*Rui Fan*



# Distributed computing

- Distributed system
  - Set of autonomous nodes, working independently of each other.
  - Nodes may be able to communicate, at a cost.
  - Ex Internet, computer cluster, sensor network.
- Nodes need to coordinate to solve some problem.
- Coordination can be done using communication. But communication is expensive.
- By making nodes randomized, they can coordinate with minimal communication.
- Randomization also simplifies symmetry breaking between nodes.
- Today we'll look at randomized contention resolution and maximal independent set.



# Contention resolution

- Set of  $n$  nodes (e.g. cellphones) want to send each other messages.
- Only one node can send at a time.
  - If two nodes send at same time, their signals interfere and both transmissions fail.
- Nodes can't communicate.
  - Communicating requires sending messages, which is the problem we're trying to solve!
  - Nodes can't coordinate to work out a schedule. They have to randomize.



# Contention resolution

- Assume system is synchronous.
  - Nodes work in rounds.
  - Each node can try to send once per round. It succeeds if and only if it's the only node to try to send in that round.
- ❖ **Algorithm** Each node tries to send with probability  $1/n$  in every round.



# Analysis

- How many rounds before all the nodes can send?
- Let  $S_{i,t}$  be the event that node  $i$  successfully sends in  $t$ 'th round.
  - $S_{i,t}$  occurs iff  $i$  tries to send in  $t$ 'th round and all other nodes do not.
- $\Pr[S_{i,t}] = 1/n \cdot (1 - 1/n)^{n-1}$ .
  - $i$  tries to send with prob.  $1/n$ , and each of  $i$ 's  $n-1$  neighbors don't send with prob.  $1 - 1/n$ .
- **Fact** For all  $n \geq 2$ ,  $1/e \leq (1 - 1/n)^{n-1} \leq 1/2$ , and  $1/4 \leq (1 - 1/n)^n \leq 1/e$ .
- So  $\Pr[S_{i,t}] \geq 1/en$ .
- $\Pr[i \text{ fails to send in } t\text{'th round}] = 1 - \Pr[S_{i,t}] \leq 1 - 1/en$ .

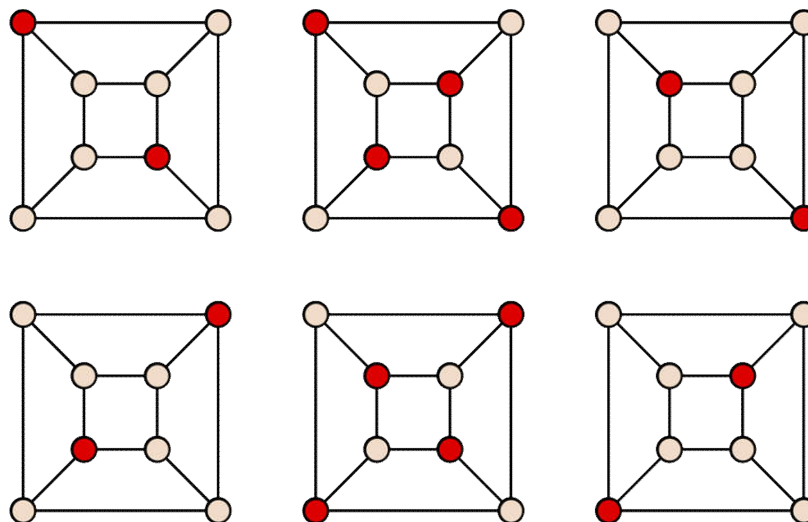


# Analysis

- **Thm** After  $2e \cdot n \ln(n)$  rounds, all nodes succeed sending with probability  $\geq 1 - 1/n$ .
- **Proof** Let  $F_i$  denote event that node  $i$  fails to send after  $2e \cdot n \ln(n)$  rounds, and let  $F$  denote event that any node fails to send after  $2e \cdot n \ln(n)$  rounds.
  - $\Pr[F_i] \leq (1 - 1/en)^{2e \cdot n \ln(n)} \leq (1/e)^{2 \ln(n)} \leq 1/n^2$ .
    - In each round  $i$  fails independently with prob.  $\leq (1 - 1/en)$ .
  - $\Pr[F] \leq \sum_i \Pr[F_i] \leq n \cdot 1/n^2 = 1/n$ , by the union bound.
  - So all nodes succeed with prob.  $\geq 1 - 1/n$ .

# Maximal independent set


- Given a graph, an independent set is a set of vertices, none of which are connected to each other.
- A maximal independent set (MIS) is an independent set such that if we add any other vertex, it would be connected to some vertex in the independent set.
  - I.e. An MIS can't be made any larger.
- A maximum independent set (MaxIS) is an independent set of maximum cardinality in the graph.
- Note that an MIS might not be a MaxIS. An MIS is a “local” max, while a MaxIS is the “global” max.



All 6 MIS's of the cube graph. Note only the two center MIS's are MaxIS.

# Distributed MIS

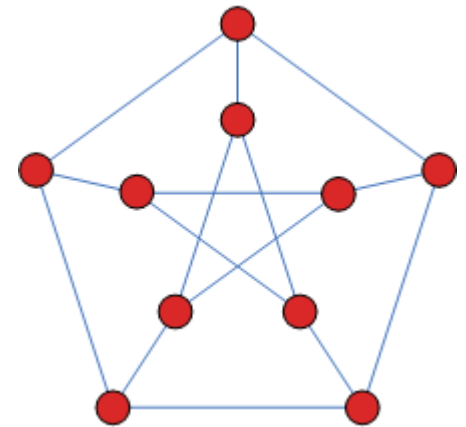
- Compute an MIS on a network of  $n$  nodes.
  - The MIS nodes can be “leaders”, used to coordinate the other nodes in some distributed computation.
- A simple algorithm is to continually add a node to the MIS, then remove its neighboring nodes and edges, then repeat.
- This algorithm takes  $O(n)$  time.
- It's also sequential. We have to remove all the neighbors of a selected node before we select the next node.
  - Otherwise we can add two neighboring nodes both to the MIS.
- We want a fast, distributed MIS algorithm.





# Distributed MIS

- Again consider synchronous model where all nodes work in rounds.
- Each node can broadcast a message to its neighbors in each round.
- ❖ Each node  $v$  chooses a random number  $r(v) \in [0, 1]$  and sends it to its neighbors.
- ❖ If  $r(v) < r(w)$  for all neighbors  $w$  of  $v$ , then  $v$  adds itself to the MIS and informs its neighbors.
- ❖ If  $v$  or one of its neighbors entered the MIS,  $v$  terminates. Remove all of  $v$ 's edges.
- ❖ Otherwise go back to first step, until graph is empty.
- Call these three steps a phase.
- Assume no ties, i.e. for any  $u, v$ , either  $r(u) < r(v)$  or vice versa.



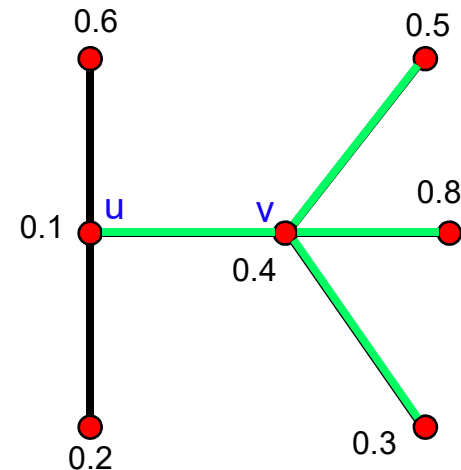


# Analysis

- We show this algorithm outputs an MIS, and terminates quickly.
- The output is an independent set.
  - For every two neighbors, only the one with smaller  $r$  value can join MIS.
  - When a node joins the MIS, all its neighbors are removed and can't join the MIS.
- It's a maximal IS because we only ever take away a node if its neighbor is in the MIS.

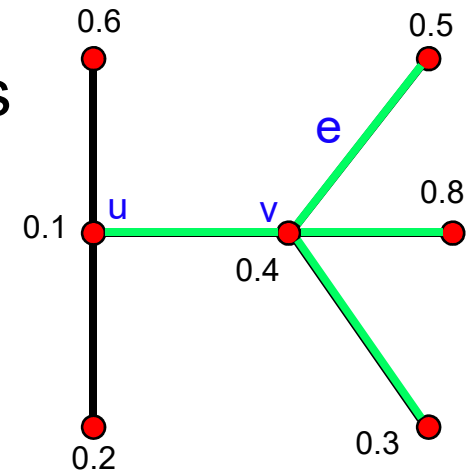
# Analysis

- How many rounds does it take to terminate?
- **Lemma** In each phase, at least half the edges are removed in expectation.
  - We'll prove this after proving Claims 1-4.
- **Def** Let  $u, v$  be two nodes. Say  $u$  preemptively removes  $v$  if  $u \in N(v)$ , and  $r(u) < r(u')$  for all  $u' \in N(u) \cup N(v)$ .
  - Denote as  $u \ll v$ .
  - Given an edge  $e = (v, w)$ , we say  $e$  is preemptively removed by  $u$  if  $u \ll v$ .
- **Claim 1** For any  $v$ , there's at most one  $u$  s.t.  $u \ll v$ .
  - **Proof** If  $u \ll v$ , then  $u$  is the neighbor of  $v$  with min  $r$  value.



# Analysis

- Let  $P = \{\text{all preemptively removed edges in phase}\}$ ,  $R = \{\text{all edges removed in phase}\}$ .
- **Claim 2**  $P \subseteq R$ 
  - Let  $e=(v,w) \in P$ . Then  $v$  has nbr  $u$  s.t.  $r(u) < r(u')$  for all  $u' \in N(u) \cup N(v)$ .
  - So  $u$  will get removed.
  - So  $v$  is also removed. All edges incident to  $v$ , including  $e$ , are also removed. So  $e \in R$ .
- Let  $X_{u < v} = 1$  if  $u < v$  and 0 otherwise.
- If  $X_{u < v} = 1$ , all edges incident to  $v$  are removed.
  - So if  $X_{u < v} = 1$ ,  $d(v)$  edges get removed, where  $d(v)$  is degree  $v$ .



# Analysis

■ **Claim 3**  $\sum_u \sum_{v \in N(u)} X_{u \ll v} * d(v) \leq 2 * |P|.$

- Given any edge  $e=(v,w)$ , the sum counts  $e$  once each time  $e$  is preemptively removed by some other node  $u$ .
- How many such  $u$ 's are there?
  - $u$  preemptively removes  $e$  only if  $u \ll v$  or  $u \ll w$ .
  - By claim 1, there's only at most one  $u$  that  $\ll v$ , and at most one that  $\ll w$ .
  - So  $e$  is preemptively removed by at most 2 other nodes.
- So any  $e$  in the sum is a preemptively removed edge that's counted at most twice.
- Since  $P$  is set of all preemptively removed edges, then the sum  $\leq 2 * |P|.$

■ **Cor 1**  $\sum_u \sum_{v \in N(u)} X_{u \ll v} * d(v) \leq 2 * |R|.$

- Because  $P \subseteq R$  by Claim 2.

# Analysis

- **Claim 4**  $E[\sum_u \sum_{v \in N(u)} X_{u < v} * d(v)] \geq |H|$ , where  $H = \{\text{edges}\}$ .
  - For any  $u$  and  $v \in N(u)$ ,  $E[X_{u < v} * d(v)] = \Pr[u < v] * d(v)$ .
  - $u < v$  only if  $r(u) < r(u')$  for all  $u' \in N(u) \cup N(v)$ .
  - There are at most  $d(u) + d(v)$  nodes in  $N(u) \cup N(v)$ .
  - Each node picks a random value  $r$ . Probability it's min among  $\leq d(u) + d(v)$  random values is  $\geq 1/(d(u) + d(v))$ .
  - So  $\Pr[u < v] \geq 1/(d(u) + d(v))$ .
  - So  $E[X_{u < v} * d(v)] \geq d(v)/(d(u) + d(v))$ .
  - $E[\sum_u \sum_{v \in N(u)} X_{u < v} * d(v)] =$   
 $\sum_{e=(u,v) \in H} (E[X_{u < v} * d(v)] + E[X_{v < u} * d(u)]) \geq$   
 $\sum_{e=(u,v) \in H} d(v)/(d(u) + d(v)) + d(u)/(d(u) + d(v)) =$   
 $\sum_{e=(u,v) \in H} 1 = |H|$ .



# Analysis

## ■ Proof of Lemma

- By Claim 4 and Cor. 1,  $|H| \leq 2 \cdot E[|P|] \leq 2 \cdot E[|R|]$ .
- So  $E[|R|] \geq |H|/2$ , i.e. half the edges get removed in expectation every phase.

## ■ Cor 2 With probability $\geq 1/3$ , at least $1/4$ the edges get removed in every phase.

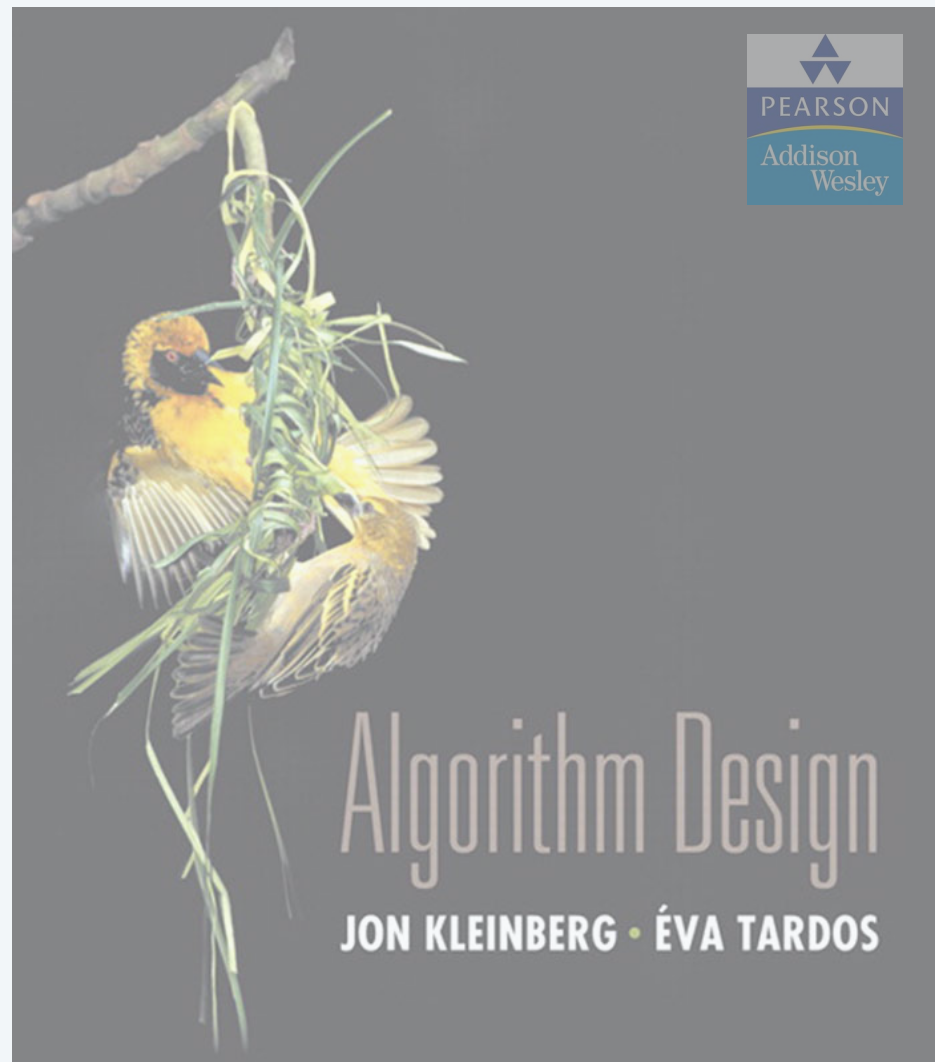
- Otherwise, the probability less than  $1/4$  edges get removed every phase is greater than  $2/3$ .
- So expected number of edges removed in the phase is  $< 2/3 \cdot |E|/4 + 1/3 \cdot |E| = |E|/2$ , contradicting the lemma.

## ■ Thm The algorithm computes an MIS in $42 \cdot \ln(n)$ phases with probability $\geq 1 - 1/n$ .

# Analysis

- **Proof** Say a phase is good if  $\geq 1/4$  the edges get removed.
  - So  $\Pr[\text{phase is good}] \geq 1/3$  by Cor 2. Also, these probabilities are independent.
  - In  $42 \cdot \ln(n)$  phases, we expect  $\geq \mu = 14 \cdot \ln(n)$  good phases.
  - $\Pr[< 7 \cdot \ln(n) \text{ good phases in } 42 \cdot \ln(n) \text{ rounds}] = \Pr[\text{number good rounds} < 1/2 \text{ expectation}] \leq e^{-14 \cdot \ln(n)/8} < 1/n$ , by Chernoff bounds.
  - If we get  $7 \cdot \ln(n)$  good phases, then fraction of remaining edges is  $\leq (3/4)^{7 \cdot \ln(n)} = n^{7 \cdot \ln(3/4)} \approx n^{-2.01}$ .
  - Since there are  $O(n^2)$  edges, all the edges get removed after  $7 \cdot \ln(n)$  good phases.
  - We get  $7 \cdot \ln(n)$  good phases in  $42 \cdot \ln(n)$  phases with probability  $> 1 - 1/n$ .





## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ ***global min cut***
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*

# Global minimum cut

---

**Global min cut.** Given a connected, undirected graph  $G = (V, E)$ , find a cut  $(A, B)$  of minimum cardinality.

**Applications.** Partitioning items in a database, identify clusters of related documents, network reliability, network design, circuit design, TSP solvers.

**Network flow solution.**

- Replace every edge  $(u, v)$  with two antiparallel edges  $(u, v)$  and  $(v, u)$ .
- Pick some vertex  $s$  and compute min  $s-v$  cut separating  $s$  from each other node  $v \in V$ .

**False intuition.** Global min-cut is harder than min  $s-t$  cut.

# Contraction algorithm

---

## Contraction algorithm. [Karger 1995]

- Pick an edge  $e = (u, v)$  uniformly at random.
- **Contract** edge  $e$ .
  - replace  $u$  and  $v$  by single new super-node  $w$
  - preserve edges, updating endpoints of  $u$  and  $v$  to  $w$
  - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes  $u_1$  and  $v_1$ .
- Return the cut (all nodes that were contracted to form  $v_1$ ).

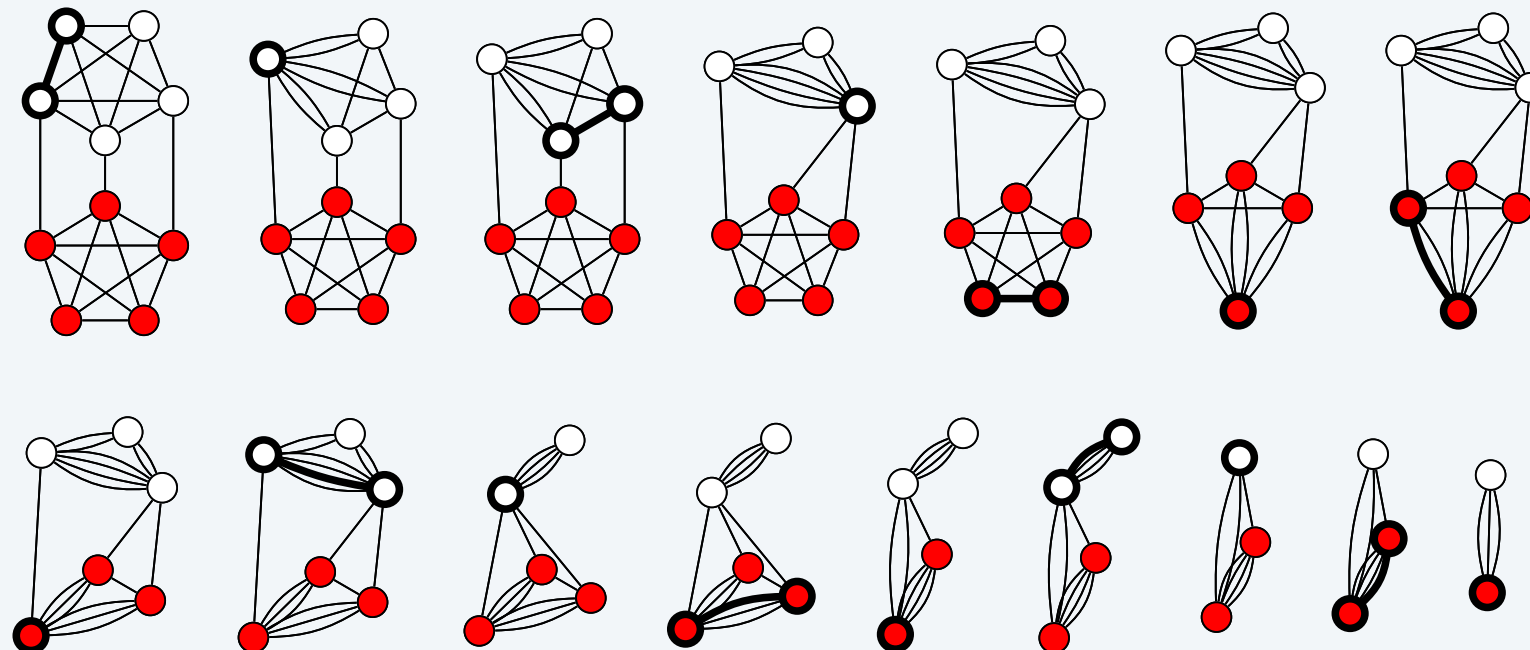


# Contraction algorithm

---

## Contraction algorithm. [Karger 1995]

- Pick an edge  $e = (u, v)$  uniformly at random.
- **Contract** edge  $e$ .
  - replace  $u$  and  $v$  by single new super-node  $w$
  - preserve edges, updating endpoints of  $u$  and  $v$  to  $w$
  - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes  $u_1$  and  $v_1$ .
- Return the cut (all nodes that were contracted to form  $v_1$ ).



Reference: Thore Husfeldt

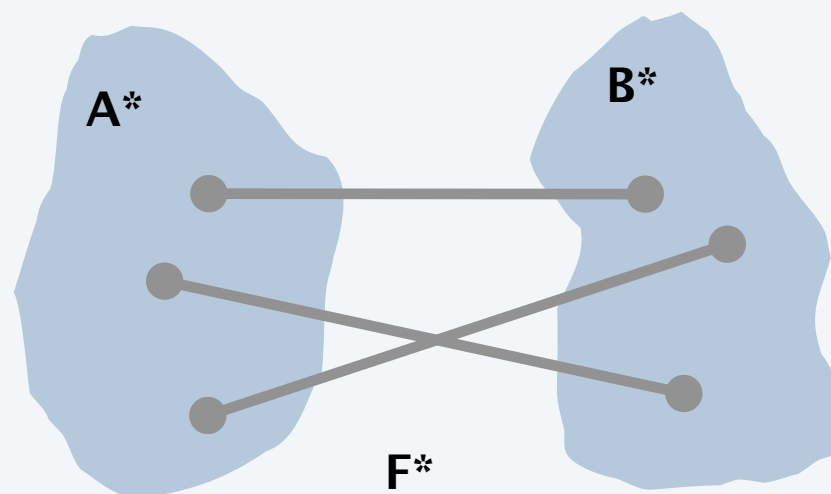
# Contraction algorithm

---

**Claim.** The contraction algorithm returns a min cut with prob  $\geq 2 / n^2$ .

**Pf.** Consider a global min-cut  $(A^*, B^*)$  of  $G$ .

- Let  $F^*$  be edges with one endpoint in  $A^*$  and the other in  $B^*$ .
- Let  $k = |F^*|$  = size of min cut.
- In **first step**, algorithm contracts an edge in  $F^*$  probability  $k / |E|$ .
- Every node has degree  $\geq k$  since otherwise  $(A^*, B^*)$  would not be a min-cut  $\Rightarrow |E| \geq \frac{1}{2} k n \Leftrightarrow k / |E| \leq 2 / n$ .
- Thus, algorithm contracts an edge in  $F^*$  with probability  $\leq 2 / n$ .



# Contraction algorithm

---

**Claim.** The contraction algorithm returns a min cut with prob  $\geq 2 / n^2$ .

**Pf.** Consider a global min-cut  $(A^*, B^*)$  of  $G$ .

- Let  $F^*$  be edges with one endpoint in  $A^*$  and the other in  $B^*$ .
- Let  $k = |F^*|$  = size of min cut.
- Let  $G'$  be graph after  $j$  iterations. There are  $n' = n - j$  supernodes.
- Suppose no edge in  $F^*$  has been contracted. The min-cut in  $G'$  is still  $k$ .
- Since value of min-cut is  $k$ ,  $|E'| \geq \frac{1}{2} k n' \Leftrightarrow k / |E'| \leq 2 / n'$ .
- Thus, algorithm contracts an edge in  $F^*$  with probability  $\leq 2 / n'$ .
- Let  $E_j$  = event that an edge in  $F^*$  is not contracted in iteration  $j$ .

$$\begin{aligned} \Pr[E_1 \cap E_2 \cdots \cap E_{n-2}] &= \Pr[E_1] \times \Pr[E_2 \mid E_1] \times \cdots \times \Pr[E_{n-2} \mid E_1 \cap E_2 \cdots \cap E_{n-3}] \\ &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \cdots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \\ &\geq \frac{2}{n^2} \end{aligned}$$

# Contraction algorithm

---

**Amplification.** To amplify the probability of success, run the contraction algorithm many times.

**Claim.** If we repeat the contraction algorithm  $n^2 \ln n$  times, then the probability of failing to find the global min-cut is  $\leq 1 / n^2$ .

with independent random choices,



**Pf.** By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^{2 \ln n} \leq \left(e^{-1}\right)^{2 \ln n} = \frac{1}{n^2}$$



$$(1 - 1/x)^x \leq 1/e$$

# Contraction algorithm: example execution



...

Reference: Thore Husfeldt



# Global min cut: context

---

**Remark.** Overall running time is slow since we perform  $\Theta(n^2 \log n)$  iterations and each takes  $\Omega(m)$  time.

**Improvement.** [Karger–Stein 1996]  $O(n^2 \log^3 n)$ .

- Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when  $n / \sqrt{2}$  nodes remain.
- Run contraction algorithm until  $n / \sqrt{2}$  nodes remain.
- Run contraction algorithm **twice** on resulting graph and return **best** of two cuts.

**Extensions.** Naturally generalizes to handle positive weights.

**Best known.** [Karger 2000]  $O(m \log^3 n)$ .



faster than best known max flow algorithm or  
deterministic global min cut algorithm